

EU AI ACT · ARTICLE 50 · N8N

The human is the gate. The ledger cannot be edited.

An n8n module that turns Article 50 transparency duties into one approval step inside any production line: label, ask a person in Slack, write an append-only record, keep a registry of every AI system on the instance.

WHY THIS EXISTS

From 2 August 2026 the deployer, not the model vendor, answers for these.

Art. 50(4) §1 · deep fakes

Generated or manipulated image, audio or video that resembles a real person, place, object or event must be labelled. Artistic or satirical work: limited disclosure.

Art. 50(4) §2 · public-interest text

AI-generated text published to inform the public must be disclosed, unless a named person took editorial responsibility.

Art. 50(1) · chatbots

A system that talks to people must make its AI nature clear unless it is obvious.

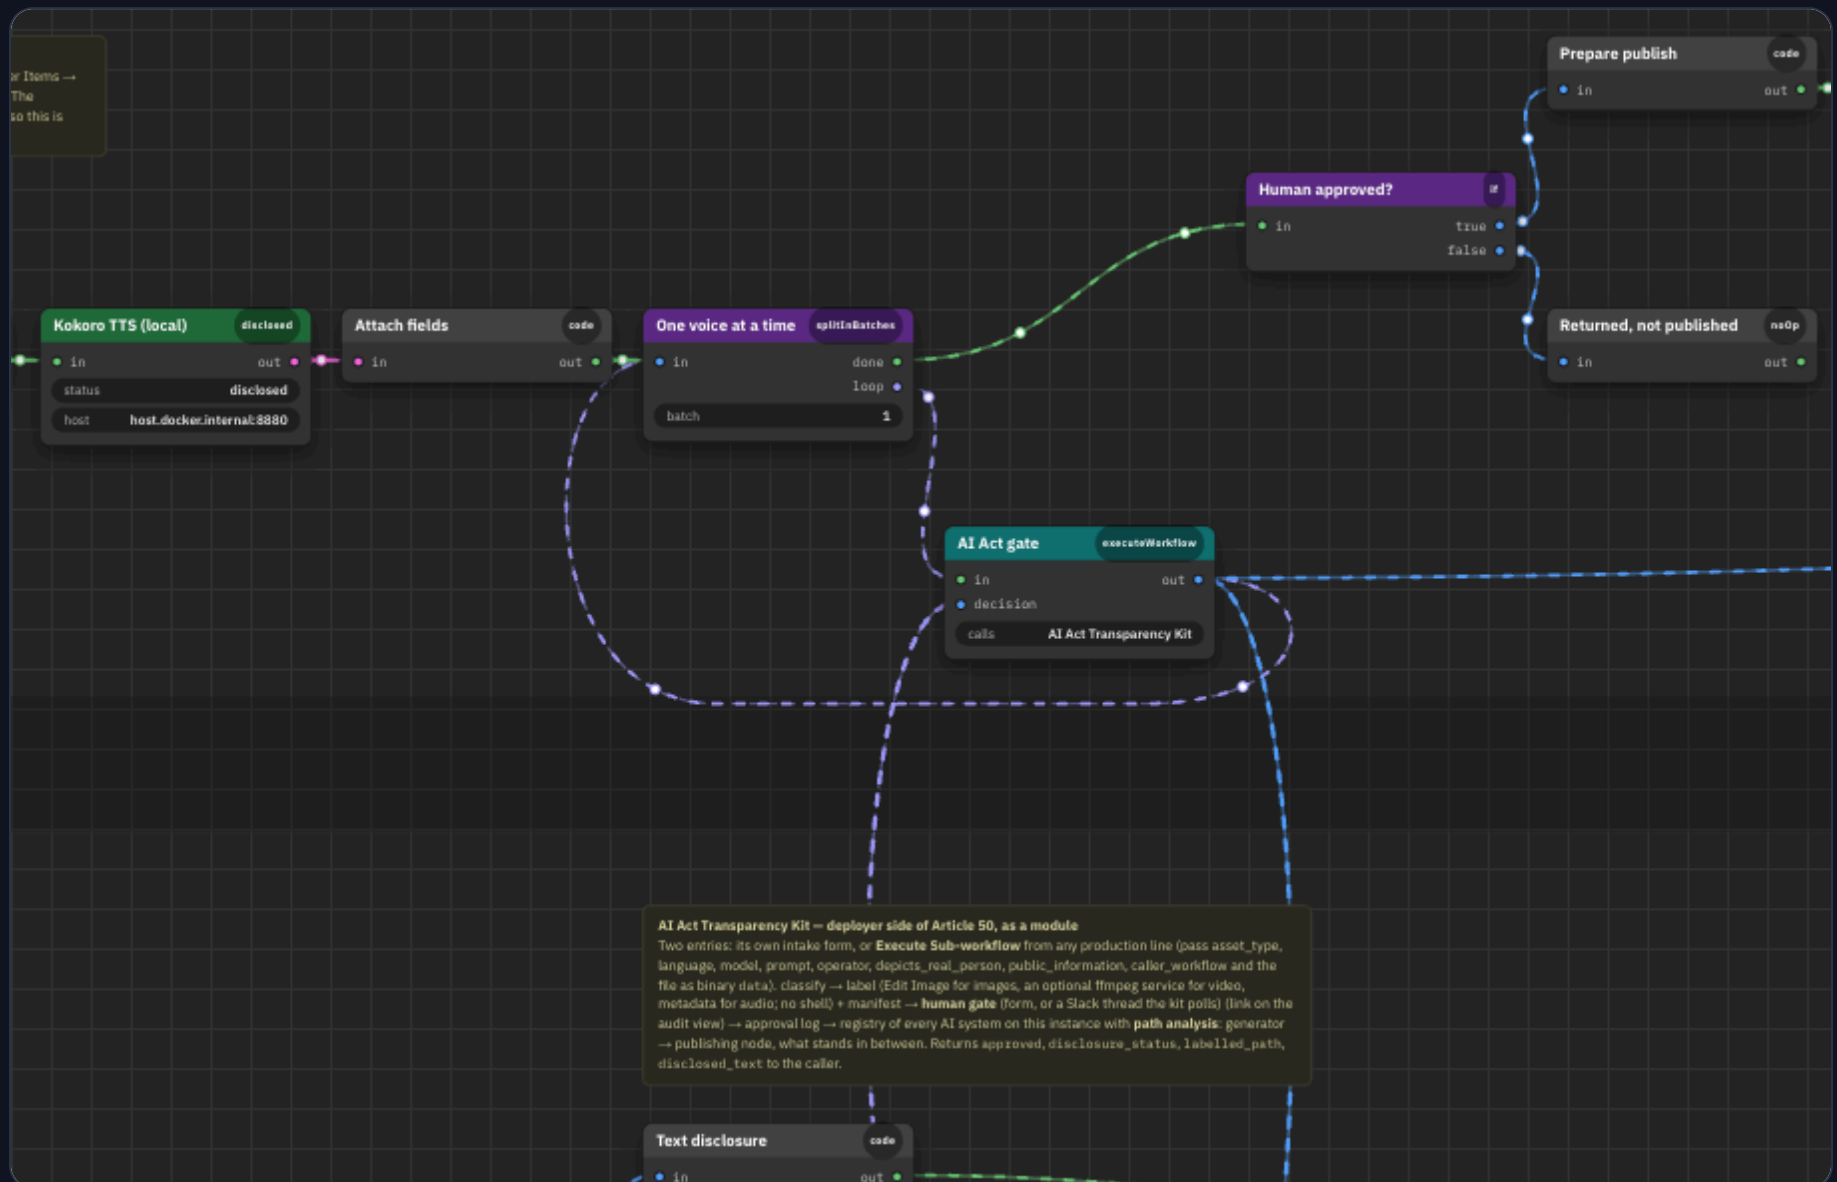
Art. 50(3) · emotion & biometrics

People exposed to emotion recognition or biometric categorisation must be told.

Most automation lines wire a generator straight into a publish node. Nothing in between knows the law. That gap is what the kit fills.

WHERE IT SITS

One sub-workflow call inside any n8n line.



A real line: notice text → three local TTS voices → **AI Act gate** → publish only what a human approved. The map is generated from the running workflows, not drawn.

- 1 The line passes the asset, the model, the prompt and what it depicts.
- 2 The kit classifies under Article 50, burns the label in, stores a manifest.
- 3 A person decides in Slack. The kit returns approved, `disclosure_status` and the labelled file.

INTAKE → LABEL

A real place, altered, is a deep fake. No person needed.



Original photo. Varna, fog.



A building added at the end of the street.

AI-manipulated label burnt in, sentence goes under the photo.

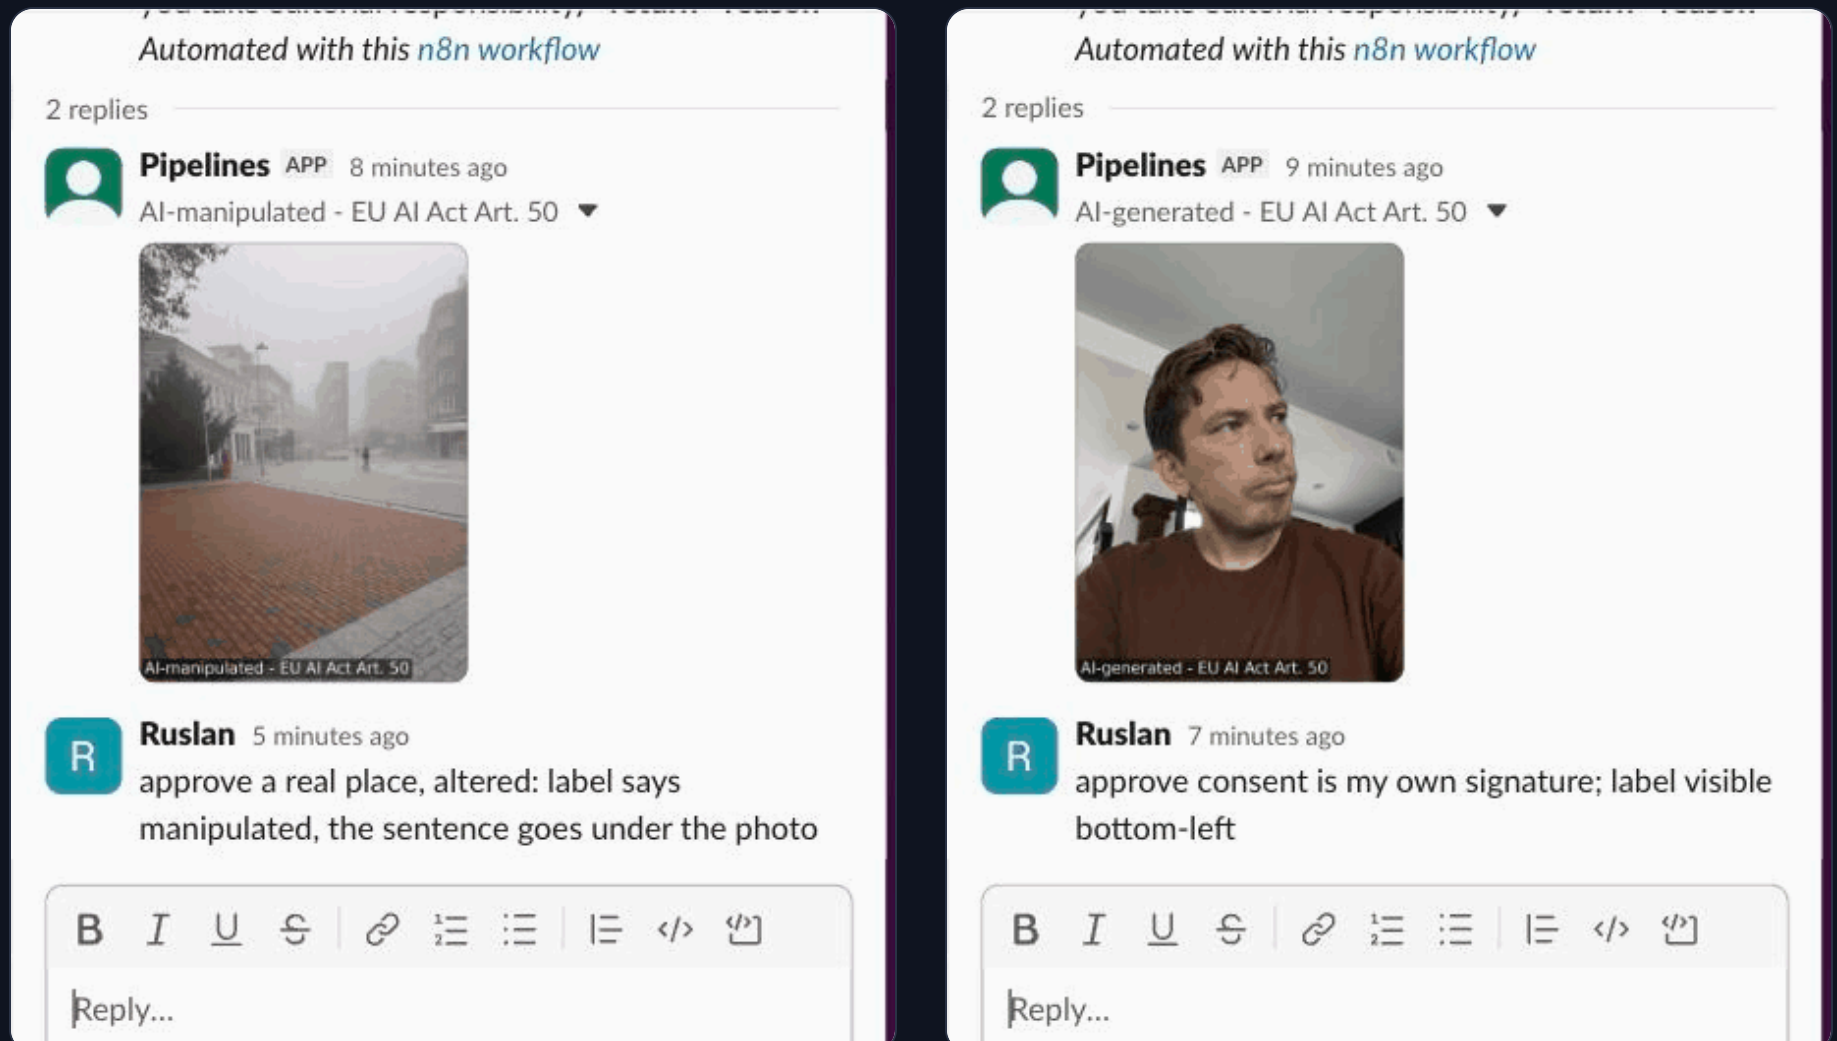


Same alteration declared as satire. Disclosure limited to a corner label, as Art. 50(4) allows.

Art. 3(60) covers persons, objects, places and events. The classifier reads four intake answers: what it shows, generated or manipulated, artistic or not, consent reference.

THE GATE

A person decides in a Slack thread. The kit only listens.



Post with the disclosure sentence and the labelled file → reply **approve**, **editorial** or **return** **<reason>**. Reviewer identity comes from Slack, not from a text field.

- Outbound polling only. n8n needs no public address, no inbound webhook.
- No answer in 24 h → the asset is returned automatically and logged as such.
- Swap Slack for a form, Teams or e-mail: the gate is one node group.

THE RECORD

Append-only ledger, hash-chained. Three real rows.

ASSET	REVIEWER	DECISION	STATUS	PREV → HASH
text · gpt-4o public notice	Ruslan Slack U0C2N5PEJJV	approve · dates checked against the municipal order	disclosed	84c55f81 → 92d911ce
image · Photoshop GF real place, manipulated	Ruslan Slack U0C2N5PEJJV	approve · label says manipulated, sentence under the photo	disclosed	5f9ee8ae → 84c55f81
image · Midjourney real person, consent on file	Ruslan Slack U0C2N5PEJJV	approve · consent is my own signature	disclosed	54c702cf → 5f9ee8ae

Postgres, two triggers.

BEFORE INSERT computes sha256 over the previous hash and the row. UPDATE and DELETE raise. `verify_chain()` walks the chain on every audit render.

What an auditor asks first.

Who decided, when, on what basis, and can the record be changed afterwards. Four answers, one table.

THE REGISTRY

Every AI system on the instance, with the path to people.

Audit view · AI Act Transparency Kit

What the auditor reads instead of the pipeline. Deployer obligations under the EU AI Act, filled automatically by the line. Generated 2026-09-17 16:23.

18

AI systems on the instance

3

people not told

2

need a decision

43

synthetic media assets

18

generated texts

61

decisions logged

0

awaiting a human

1

AI-systems registry

voluntary · supports Art. 4 literacy

Every AI system on this n8n instance, with path analysis: from each node that calls a model, every path forward to a node that reaches people, and what stands in between. Plus every external generator declared at intake. Read from the instance, not typed by hand.

STATUS	WORKFLOW / LINE	SYSTEM (NODE)	NODE TYPE	MODEL	REACHES PEOPLE VIA	EVIDENCE (PATH)	SOURCE	ASSETS
uncovered	Sample: social post drafter (no gate) inactive	Draft post with GPT	langchain.openAi	gpt-4o-mini	Publish to LinkedIn	Draft post with GPT → Voice-over (ElevenLabs) → Publish to LinkedIn — no disclosure, no human gate on this path	workflow node	
uncovered	Sample: social post drafter (no gate) inactive	Voice-over (ElevenLabs) face/voice	httpRequest	api.elevenlabs.io	Publish to LinkedIn	Voice-over (ElevenLabs) → Publish to LinkedIn — no disclosure, no human gate on this path	workflow node	
inform	Sample: support inbox mood scoring (no notice) inactive	Score customer emotion (Hume) emotion/biometric	httpRequest	api.hume.ai		emotion recognition or biometric categorisation runs on people — they must be informed (Art. 50(3)); no node tagged [persons informed] in this workflow	workflow node	
chatbot	Sample: website chatbot (no AI notice) inactive	Visitor writes in chat talks to people	langchain.chatTrigger			interacts with natural persons — the provider must make the AI nature clear unless obvious (Art. 50(1)); no node tagged [ai disclosed] in this workflow	workflow node	
chatbot	Sample: website chatbot (no AI notice) inactive	Answer as 'Anna from support' talks to people	langchain.agent			interacts with natural persons — the provider must make the AI nature clear unless obvious (Art. 50(1)); no node tagged [ai disclosed] in this workflow	workflow node	
disclosed	Recycling notice · three voices active	Kokoro TTS (local)	httpRequest	host.docker.internal:8880	Publish to site [exit]	Kokoro TTS (local) → Attach fields → One voice at a time → Human approved? → Prepare publish → Read approved file → Publish to site [exit]	workflow node	

The kit reads all workflows through the n8n API and walks from every model node forward. What stands between the generator and the exit decides the status:

THE AUDIT VIEW

What the auditor reads instead of the pipeline.

2 Synthetic media and deep fakes

Images, video and audio. Deep fake per Art. 3(60): generated or manipulated content resembling a real person, object, place or event that would appear authentic if not for the fact that it is generated or manipulated. Provenance manifest for all of them.

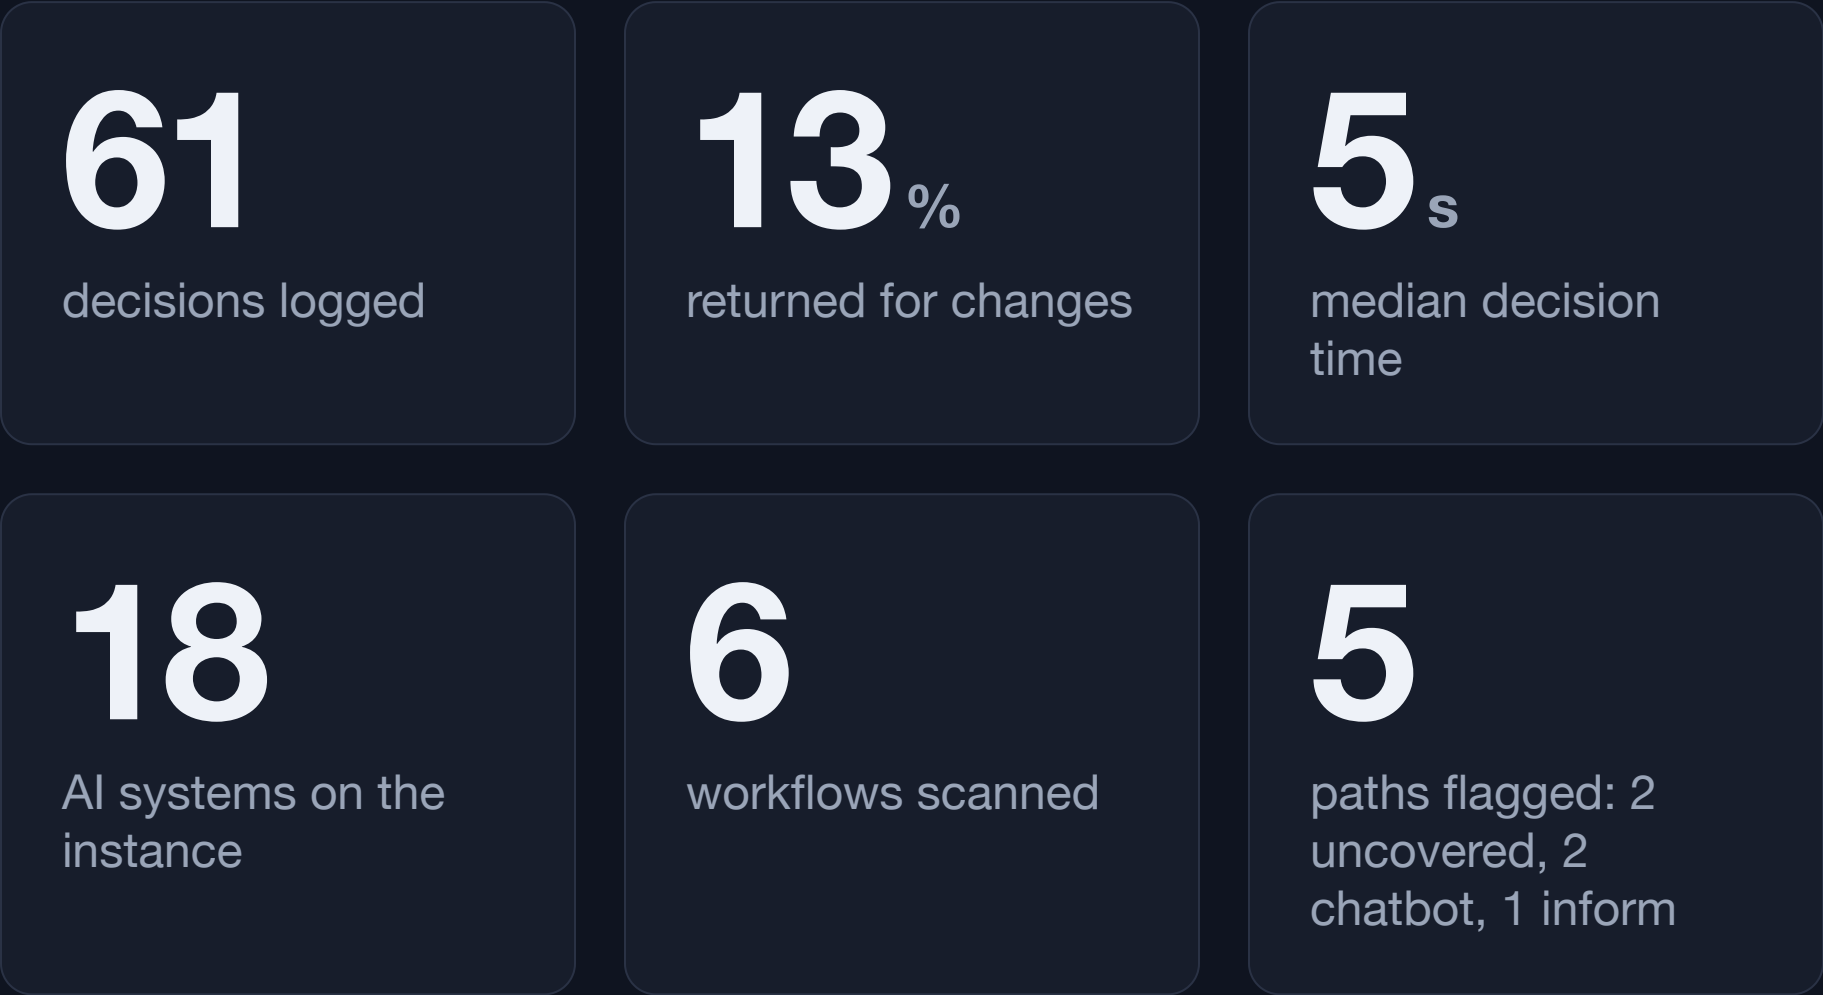
ID	TYPE	CATEGORY	SHOWS	CONSENT	DISCLOSURE	MODEL	PROMPT HASH	OPERATOR	STATUS
2026-09-17-16-11-35-323-d0750b817r	image	deepfake	a real person (face or voice)	consent/2026-09-16-rk-likeness.pdf	burnt-in label (law)	Midjourney v7	d0750b416fc1	Ruslan Karymov	OK
2026-09-17-16-11-32-048-ca44f84wl2	image	deepfake	a real object, place or event · manipulated		burnt-in label (law)	Photoshop Generative Fill 2026	ca44f83086c1	Ruslan Karymov	OK
2026-09-17-16-01-22-129-c77a24wije	image	deepfake	a real object, place or event · manipulated		burnt-in label (law)	Photoshop Generative Fill 2026	c77a242cfba5	Ruslan Karymov	OK
2026-09-17-13-38-01-241-611a0caum8	video	synthetic_media	nothing real		burnt-in label (house policy)	Runway Gen-4	611a0ce28bdd	Ruslan Karymov	OK
2026-09-17-13-37-55-937-472390kcmo	audio	deepfake	a real person (face or voice)	consent/2026-09-16-rk-voice.pdf	metadata + sentence at publication (law)	ElevenLabs Multilingual v2	4723904f3e3a	Ruslan Karymov	OK
2026-09-17-13-37-50-608-468c16n6pz	image	deepfake	a real object, place or event · manipulated · artistic		burnt-in label (law (limited: artistic work))	Photoshop Generative Fill 2026	468c165a27bd	Ruslan Karymov	OK

Section 2 of the audit view: every synthetic asset with what it shows, consent reference, disclosure method, model, prompt hash, operator, status, artefact path.

- One GET endpoint, rendered from the ledger. No spreadsheet kept by hand.
- Four sections: registry (Art. 4 literacy, voluntary) · synthetic media (Art. 50(4) §1) · generated text (Art. 50(4) §2) · approval log.
- Legal basis on every block. Voluntary items are labelled voluntary.

NUMBERS FROM THE LEDGER

Demo instance, one day, one reviewer.



These are demo numbers, not a production benchmark: I answered my own requests within seconds. What the numbers prove is the mechanics: every decision landed, the chain verified after each run, nothing was edited by hand.

TAKE IT

Import the JSON. Point it at your line. Add a Slack channel.

Repository

github.com/karusrus/opacity-kit

Case page and live map

karusrus.github.io/opacity-kit ·

karusrus.github.io/pipeline-map

What I do in a team's first month: inventory every line that calls a model, put the gate where content reaches people, hand compliance a page they can read without opening n8n.

Ruslan Karymov

AI & Automation Lead · open to EU remote and Bulgaria

linkedin.com/in/ruslankarymov